

MASTER PROJECT

Shock propagation through the lens of remittances

Who bears the cost? Evidence from US–Mexico migration corridors

Authors

Anqi Wang Liu, Francesco Zucca, Joan Costa Rodríguez,

Michael Fehl, Xènia Alemany Rovira

Barcelona School of Economics

Master's program in Economics

Contents

Abstract	ii
1 Introduction	1
2 Theoretical framework and empirical strategy	5
2.1 A Remittance Model with Network Spillovers and Heterogeneous Migrants . .	5
3 Empirical Context and Data	9
3.1 The Natural Experiment: Hurricane Ian (2022)	9
3.2 Preliminary statistics	9
4 Data and Bilateral Remittance Flow Estimation	10
4.1 Data Sources	10
4.2 Migration Networks Construction	11
4.3 Underlying assumptions and validation exercises	11
4.3.1 Persistence of migration networks	12
4.3.2 Validation of the estimated remittance flows	13
5 Empirical Methodology and Results	14
5.1 Methodology	14
5.2 Robustness checks	14
5.3 Results	15
6 Investigating Spillover Effects	16
6.1 Empirical Strategy	16
6.2 Results	18
7 Conclusion	19
7.1 Summary of Findings	19
7.2 Policy Implications	19
7.3 Limitations and Future Research	19
Bibliography	20

Abstract

1. Introduction

International remittances are one of the main private channels through which local economic shocks can cross national borders. Mexico is a central setting in which to study this mechanism. Banco de México records US\$64.7 billion in family remittance receipts in 2024, while World Bank estimates place Mexico among the largest remittance recipients in the world, second only to India (Banco de México, 2026b; World Bank, 2024). The U.S.–Mexico corridor is especially important because Mexican households receive transfers from migrants whose earnings are tied to local labor markets across the United States. This paper asks whether a negative local shock in the United States is transmitted to Mexican municipalities through remittances. Do migrants smooth transfers after being hit by a shock, or do origin households bear part of the adjustment through lower remittance inflows?

The answer matters because remittances are not only an aggregate balance-of-payments flow, they are also part of household resources. Our data do not directly observe Mexican household income or consumption. Even so, changes in remittance inflows are economically meaningful for household welfare. Prior work shows that remittances are an important source of household resources in Mexico, affecting consumption, poverty, and inequality (Taylor et al., 2005), food security (Mora-Rivera & van Gameren, 2021), and schooling and child labor decisions (Alcaraz et al., 2012). We therefore interpret changes in municipal remittance inflows as reduced-form evidence on a channel with direct relevance for household income and consumption, while not claiming to estimate those household outcomes themselves.

The main empirical difficulty is that most large shocks are not cleanly assigned to one side of the migration corridor. Caballero et al. (2023) show that Mexican households connected to U.S. labor markets more severely affected by the Great Recession were less likely to receive remittances, providing important evidence that migrant networks transmit U.S. shocks to Mexico. Yet the Great Recession simultaneously affected the United States and Mexico, making it difficult to isolate the remittance channel from other aggregate forces. Their analysis also focuses on the extensive margin of remittance receipt. Quantifying how much bilateral remittance flows adjust on the intensive margin is necessary to measure the income loss experienced by origin households.

We address these challenges by exploiting Hurricane Ian, which made landfall in southwestern

Florida on September 28, 2022. Ian caused an estimated US\$109.5 billion in damages in Florida and was the costliest hurricane in the state's history (Bucci et al., 2026). The shock was severe and spatially concentrated on the sender side of the corridor: it disrupted economic activity in Florida but did not directly affect Mexican municipalities. Natural disasters in the United States can reduce local productivity, labor demand, housing wealth, and household resources (Boustan et al., 2020). Hurricane Ian therefore provides a localized negative shock to the capacity of Florida-based migrants to remit, while leaving recipient municipalities in Mexico exposed only through their pre-existing migration links to Florida.

We base our argument on the following idea. A shock to Florida migrants can directly reduce remittances from Florida to Mexican municipalities. However, many origin municipalities receive remittances from migrants in several U.S. states. If migrants care about total resources available to the origin household, then a fall in Florida remittances can raise the marginal value of transfers from migrants living elsewhere. The same shock can therefore generate two effects at once: a direct decline in Florida-to-Mexico remittances and a compensatory response from other U.S. states connected to the same municipalities.

The first part of the paper formalizes this idea with a structural remittance model. The model follows the insight in Albert and Monras (2022) that migrants allocate resources across where they live and where their origin households are located. In our setting, this idea is applied to remittance choices rather than to location choices. Migrants choose between local consumption in the United States and resources available to the origin household in Mexico. They differ in income and in the strength of their altruistic or obligation motive, and remitting involves both a fixed cost and a variable cost. In the spirit of models with heterogeneous agents, fixed costs, and variable costs (Melitz, 2003), the fixed cost generates an extensive margin of remitting, while the variable cost governs the intensity of transfers.

The key extension is that origin-household resources include local non-remittance income plus remittances received from migrants in other U.S. states. This makes the model directly useful for interpreting the empirical patterns. A negative shock to Florida migrants reduces their optimal remittances. At the same time, for households linked to Florida, the fall in Florida transfers lowers perceived origin resources for migrants in other states, raising the marginal utility of additional remittances under an insurance motive. The model therefore provides a unified way to interpret both the direct exposure effect and the compensatory network re-

sponse observed in the data.

The second part of the paper brings the model to the data by constructing the inflow of remittances flows from U.S. states to Mexican municipalities. Because such flows are not directly observed, we combine two sources from Banco de México with MCAS administrative records. Banco de México reports remittance inflows by U.S. state of origin and by Mexican municipality of receipt (Banco de México, 2026a, 2026c). MCAS records identify the Mexican municipality of origin and U.S. state of residence of Mexican consular-card holders, providing a proxy for the geography of migrant networks. We use these records to construct state-to-municipality migration shares and allocate observed state-level remittance outflows across Mexican municipalities. We focus our analysis on the Mexican municipalities receiving remittances from Florida.

Two validation exercises support this construction. First, the migration network data are highly persistent: the distribution of Mexican migrants across U.S. states and Mexican municipalities has year-to-year correlations above 0.90 from 2012 onward. Second, when we use these migration shares to allocate state-level remittance outflows to Mexican municipalities, the resulting municipal aggregates correlate between 0.75 and 0.85 with official municipal remittance receipts. These facts suggest that the estimated corridors capture meaningful variation in the geography of remittances.

The third part of the paper studies the shock. We exploit variation in Mexican municipalities' pre-existing exposure to Florida to estimate whether Ian affected municipal remittance inflows differently depending on their exposure to the shocked corridor. In the period immediately following Ian, a 1% increase in a municipality's pre-existing exposure to Florida is associated with a 0.16% decrease in remittances received. This result is consistent with a sender-side income shock being transmitted through the remittance corridor to origin municipalities.

Finally, we examine migration-network spillovers. Among municipalities jointly exposed to Florida and to other large sender states, remittances from California are insignificant but decrease in Texas relative to Florida in the subsequent post-Ian quarters. We explore this relationship further by centering the data on different quantiles of Texas exposure to Mexican municipalities. The diagnostic plots show for the 10% most exposed Florida municipalities to Texas, (insert write-up here), whereas for the least 10% exposed show (describe what we see). This pattern is consistent with both compensatory transfers from migrants in other states and

out-migration from Florida-exposed counties.

Related Literature. The paper is closest to work that studies remittances as a form of household insurance (Amuedo-Dorantes & Pozo, 2006; Bettin et al., 2025; Yang & Choi, 2007). Most existing evidence studies shocks in receiving countries, where migrants abroad may increase transfers to insure households at origin. We instead study a sender-side shock, where the capacity to remit changes in the destination labor market while origin municipalities are not directly affected. This distinction is important because the same remittance channel that insures households against origin shocks may transmit destination shocks back to those households.

The paper also relates to research on migrant networks and the international transmission of local shocks (Caballero et al., 2023; Munshi, 2003). The closest paper is Caballero et al. (2023), which shows that Mexican households linked to U.S. labor markets hit by the Great Recession became less likely to receive remittances. We differ by studying a localized natural disaster, constructing bilateral state-to-municipality remittance corridors, and focusing on intensive-margin flows and cross-corridor responses.

The structural framework connects to research that models migrants as allocating resources across origin and destination locations (Albert & Monras, 2022). The difference is that our model uses this logic to study remittance behavior after a destination-side shock. It also connects to work on interference and spillovers in causal inference (Aronow & Samii, 2017; Butts, 2023; Hudgens & Halloran, 2008). In our setting, interference is not only a threat to identification. It is part of the economic mechanism, because migrants linked to the same origin municipality may jointly insure households against shocks that affect only one destination.

The remainder of the paper is organized as follows. Chapter 2 develops the theoretical remittance model. Chapter 3 introduces Hurricane Ian as the natural experiment and describes both the construction and validation of the migration and remittance data. Chapter 4 presents the empirical strategy and reports the main estimates for exposure effects. Chapter 5 investigates potential spillover effects. Chapter 6 concludes.

2. Theoretical framework and empirical strategy

This chapter outlines the theoretical framework used for our analysis. We begin by adapting a heterogeneous-firm trade model to the context of individual remittance behavior, establishing the theoretical foundations of how a migrant's transfer decisions can be influenced by economic shocks.

2.1. A Remittance Model with Network Spillovers and Heterogeneous Migrants

To analyze the impact of shocks on remittance flows, we adapt the logic of heterogeneous-firm trade models to the context of remittances. Specifically, we consider a model where migrants differ in income and altruism intensity, and where remitting requires paying both a fixed access cost and a variable cost proportional to the amount remitted. We then augment the model to allow migrants connected to the same origin municipality to respond to each other's transfers. Formally, consider migrant i , who resides in U.S. state s , originates from Mexican municipality m , and chooses remittances in period t .

We decompose origin-household resources into local non-remittance income and remittances received from migrants in other U.S. states. Perceived origin-household resources excluding migrant i 's own remittance are

$$y_{imst}^O = \tilde{y}_{mt}^O + \sum_{s' \neq s} R_{s'mt}, \quad (2.1)$$

where $\tilde{y}_{mt}^O$ denotes local non-remittance income in municipality m , and $\sum_{s' \neq s} R_{s'mt}$ captures remittance inflows from migrants residing in U.S. states other than s . The migrant faces a utility maximization problem under a CES aggregator over local consumption and total resources available to the origin household:

$$U_{imst} = \left[\alpha \left(c_{imst}^L \right)^{\frac{\sigma-1}{\sigma}} + (1-\alpha) \left(\phi_{im} \left[R_{imst} + \tilde{y}_{mt}^O + \sum_{s' \neq s} R_{s'mt} \right] \right)^{\frac{\sigma-1}{\sigma}} \right]^{\frac{\sigma}{\sigma-1}}. \quad (2.2)$$

Here $\phi_{im} > 0$ denotes the altruism or obligation intensity parameter, analogous to the productivity parameter in Melitz (2003). This strictly positive parameter ensures that sending remit-

tances can increase the migrant's utility by raising origin-household resources. The variable $c_{imst}^L \geq 0$ represents the migrant's consumption in the destination country, and $R_{imst} \geq 0$ is the amount sent in remittances by migrant i . The weight of local consumption is given by $\alpha \in (0, 1)$, and the elasticity of substitution between local consumption and origin-household resources is given by σ . Origin-household total resources are therefore defined as $R_{imst} + y_{imst}^O$, which we treat as a single composite good. The presence of y_{imst}^O introduces a role for both origin-country economic conditions and migration-network transfers: lower perceived origin resources increase the marginal utility of remittances, generating an insurance motive consistent with Yang and Choi (2007).

The migrant's budget constraint is given by:

$$c_{imst}^L + f_{mt} \mathbb{1}\{R_{imst} > 0\} + (1 + \tau_{mt})R_{imst} \leq w_{imst}, \quad (2.3)$$

where $w_{imst} \geq 0$ is the migrant's wage, $f_{mt} \geq 0$ is a fixed remitting cost capturing the per-transaction expense of accessing a remittance channel, analogous to the fixed export cost in Melitz (2003), and $\tau_{mt} > -1$ is the variable cost which scales with the amount remitted, analogous to the iceberg cost structure in Melitz (2003). Notice that, even though τ_{mt} is typically positive to reflect transaction fees or exchange rate spreads, we allow for the possibility of negative variable costs, which could arise from the presence of remittance subsidies or favorable exchange rate regimes.

Assuming that the migrant maximizes utility subject to the budget constraint and chooses to remit a positive amount ($R_{imst} > 0$), solving the Lagrangian yields the optimal remittance function:

$$R_{imst}^* = \frac{A_{imst} [w_{imst} - f_{mt} + (1 + \tau_{mt}) (\tilde{y}_{mt}^O + \sum_{s' \neq s} R_{s'mt})]}{1 + (1 + \tau_{mt})A_{imst}} - \left(\tilde{y}_{mt}^O + \sum_{s' \neq s} R_{s'mt} \right), \quad (2.4)$$

where we define the parameter A_{imst} as:

$$A_{imst} \equiv \left(\frac{(1 - \alpha)\phi_{im}}{\alpha(1 + \tau_{mt})} \right)^\sigma. \quad (2.5)$$

The term A_{imst} represents the optimal ratio of origin-household resources to own local consumption, adjusted for the variable cost of remitting. It comprises the migrant's altruism in-

tensity (ϕ_{im}), the utility weight of own consumption (α), the elasticity of substitution (σ), and the variable cost of remitting (τ_{mt}).

To understand the effect of destination-country economic shocks on remittance flows, we analyze changes in the migrant's wage (w_{imst}). The partial derivative of optimal remittances with respect to wage is strictly positive:

$$\frac{\partial R_{imst}^*}{\partial w_{imst}} = \frac{A_{imst}}{1 + (1 + \tau_{mt})A_{imst}} > 0. \quad (2.6)$$

Therefore, through an income effect, the model predicts that a positive shock to the migrant's wage will increase remittances, while a negative shock will decrease them. The fixed cost f_{mt} establishes an entry threshold: if the migrant's wage is not sufficiently high relative to the fixed cost, they will not send remittances at all. Additionally, the variable cost τ_{mt} reduces the marginal benefit of each dollar sent, dampening the sensitivity of remittances to wage fluctuations.

The augmented model also delivers a spillover prediction. Holding the migrant's own wage and remitting costs fixed, remittances from other migrants enter the decision through perceived origin-household resources:

$$\frac{\partial R_{imst}^*}{\partial R_{s'mt}} = -\frac{1}{1 + (1 + \tau_{mt})A_{imst}} < 0 \quad \text{for } s' \neq s. \quad (2.7)$$

A fall in remittances from one destination lowers perceived resources at origin and raises the optimal remittance from migrants in other destinations. For municipalities highly exposed to the shocked destination, this lowers origin-household resources and increases the marginal utility of additional transfers from migrants residing elsewhere. The model therefore provides a theoretical basis for compensatory migration-network spillovers: the same shock can generate a direct fall in remittances from the affected destination and an offsetting response from other connected destinations.

The model delivers three empirical implications that guide the analysis. First, because optimal remittances are increasing in migrant income, a negative wage or income shock in the country where migrants reside should reduce remittances sent by affected migrants. Second, origin municipalities with greater pre-existing exposure to the affected migrant destination should experience larger declines in total remittance receipts, all else equal. This exposure effect is a

reduced-form object: it captures the direct loss in transfers from the affected destination net of any adjustment by migrants in other destinations. Third, because remittances from other destinations enter the migrant's decision through perceived origin-household resources, a decline in transfers from the affected destination raises the marginal value of remittances from migrants in alternative destinations connected to the same municipality. The model therefore predicts compensatory responses from other high-sending destinations, especially for municipalities jointly exposed to the affected destination and those alternatives. These implications map directly into the empirical strategy below: we first estimate whether remittances fall from the shocked destination, then test whether municipal receipts fall more in places more exposed to that destination, and finally examine whether alternative corridors increase transfers relative to the shocked corridor among jointly exposed municipalities.

3. Empirical Context and Data

3.1. The Natural Experiment: Hurricane Ian (2022)

The shock selected for this analysis is Hurricane Ian, which struck the southwestern coast of Florida on September 28, 2022, with peak sustained winds of 140 kt (around 260 kph) (Bucci et al., 2026). Ian caused an estimated US\$ 112.9 billion in total damage in the United States, of which US\$ 109.5 billion occurred in Florida alone, making it simultaneously the costliest hurricane in Florida's history and the third-costliest in US history. In Lee County alone, at least 52,514 structures were impacted, of which 5,369 were completely destroyed (Bucci et al., 2026).

Crucially for our identification strategy, Hurricane Ian did not affect Mexico. The storm originated in the Caribbean, made landfall in Cuba and Florida, and dissipated over the Atlantic after a final less impactful landfall in South Carolina (Bucci et al., 2026). Therefore, the shock is spatially clean. Indeed, remittance senders in Florida were directly exposed to a severe and sudden wealth shock (reference or data?), while remittance receivers in Mexican municipalities were entirely unaffected. This allows us to better isolate the transmission of the shock through the remittance channel. In the language of the model, Hurricane Ian is the negative shock to migrants' economic resources in the country where they reside, Florida is the affected destination, and Mexican municipalities differ in their pre-existing exposure to Florida through migration networks.

****miguel to do: INCLUDE WRITE UP AND ANALYSIS ABOUT IT AFFECTING MEXICAN-BORN WORKERS IN FLORIDA AFFECTED SECTORS****

3.2. Preliminary statistics

****insert Xenia****

4. Data and Bilateral Remittance Flow Estimation

This chapter outlines the data sources and the methodology we construct for estimating bilateral remittance flows between U.S. states and Mexican municipalities, which is not publicly available. First, we describe the datasets from Banco de México and Matricula Consular de Alta Seguridad (MCAS), which provide the remittance and migration data, respectively. Then, we compute a migration-share weighting matrix, which we use to distribute U.S. state-level remittance outflows to Mexican municipalities. Finally, we explore the underlying assumptions of our estimation approach, testing the persistence of migration networks over time and validating our estimates by comparing them with official data available at a more aggregate level.

4.1. Data Sources

The source of remittance data is Banco de México (Banxico), which provides two complementary datasets through its Sistema de Información Económica (SIE). The first, series CE168, records remittance inflows from each US state to Mexico in millions of current US dollars at quarterly frequency, covering the period from Q1 2013 to Q1 2026 (Banco de México, 2026a). This series serves as the US state outflows (r_s) for our bilateral flow estimation. The second, series CE166, records remittance inflows received by each Mexican municipality, also in millions of current US dollars at quarterly frequency and over the same sample period (Banco de México, 2026c). This series acts as a benchmark to validate the accuracy of our estimates at an aggregate level.

To bridge these datasets, we utilize administrative records from the Matricula Consular de Alta Seguridad (MCAS) program, managed by the Instituto de los Mexicanos en el Exterior (IME) (Instituto de los Mexicanos en el Exterior, 2026). The publicly available version of this dataset provides a comprehensive registry of consular identification cards issued to Mexican nationals residing in the United States each year, from 2010 to 2024. While the MCAS does not constitute an exhaustive census of all migrants, it represents a robust proxy for the spatial distribution of Mexican migrants in the United States. In particular, it is highly representative of undocumented and recently arrived migrants, who face the strongest incentives to obtain consular identification (Caballero et al., 2018).

4.2. Migration Networks Construction

To estimate bilaterance remittance flows from U.S. states to Mexican municipalities, we construct a migration-based weighting matrix using annual MCAS from 2010 to 2024. Because these records capture new consular issuances rather than the complete stock of the remitting population, which also includes previously arrived migrants, we rely on the assumption of persistence of migration patterns over time. To mitigate the volatility of yearly migration fluctuation and construct a robust proxy for the long-term spatial distribution of Mexican migrants, we compute the average of these matrices across all years. We then apply this averaged weighting matrix to the state-level remittance outflows from Banxico's series CE168 to derive the U.S. to Mexican municipality bilateral flows.

Formally, we define $MCAS_{ms}$ as the average annual count of MCAS issuances to migrants from Mexican municipality m residing in U.S. state s . For each U.S. state, we calculate the average migration share, w_{ms} , which represents the proportion of that state s 's Mexican population originating from the m municipality, as the ratio of $MCAS_{ms}$ to the total number of MCAS issuances from that state across all municipalities M :

$$w_{ms} = \frac{MCAS_{ms}}{\sum_{m=1}^M MCAS_{ms}} \quad (4.1)$$

By applying the weights computed from equation (4.1) to the official aggregate state-level outflows from Banxico (r_s), we derive the estimated bilateral remittance flow ($\hat{r}_{ms}$) from state s to municipality m :

$$\hat{r}_{ms} = r_s \cdot w_{ms} \quad (4.2)$$

This migration share estimation approach yields the final estimate of our outcome variable.

4.3. Underlying assumptions and validation exercises

The validity of this estimation relies on two primary identification assumptions. First, we assume the temporal persistence of migration flows, implying that the distribution of migrants recorded in the 2010–2024 MCAS data is representative of the active remitting population. Second, we assume that, within a given U.S. state, the average amount remitted does not vary

systematically by the migrant's municipality of origin. While this assumption ignores micro-level heterogeneity, it provides a baseline estimate of the bilateral remittance flows that can be validated against the aggregate data provided by Banxico' series CE166.

4.3.1. *Persistence of migration networks*

The persistence assumption of migration networks is critical to our estimation strategy, as it justifies the use of MCAS data to construct the migration-based weighting matrix in order to allocate remittance flows. The stability of migration patterns over time is well-documented in the literature, with studies showing that established migration networks tend to dictate future migration patterns, creating persistent flows between specific origin and destination pairs (Card, 2001; Munshi, 2003).

To verify this assumption within our specific sample period, we conduct a validation exercise using the MCAS data by computing the Pearson correlation of our constructed weighting matrices across years. As shown in Figure 4.1, when taking any year from 2012 onward as the base year, the correlation with subsequent years is consistently above 0.90. While the 2010 and 2011 matrices exhibit strong correlation between them, their correlation with later years is smaller, specifically between 0.75 and 0.85. This may be due to structural adjustment in migration patterns following the Great Recession.




Figure 4.1: Correlation of MCAS-based weighting matrices across years. Each point represents the correlation coefficient between the weighting matrix of the base year and the weighting matrix of the comparison year.

Source: Authors' own elaboration based on MCAS data.

This high degree of correlation suggests that the spatial distribution of Mexican migrants across U.S. states has remained highly stable over time, lending credibility to our use of an averaged weighting matrix for estimating bilateral remittance flows.

4.3.2. *Validation of the estimated remittance flows*

Our methodology to estimate bilateral remittance flows implicitly assumes a homogeneous send behavior among migrants within a given U.S. destination. That is, migrants in U.S. states send on average the same amount of remittances, regardless of their specific Mexican municipality of origin m . To evaluate the empirical validity of this assumption, we aggregate our estimated bilateral municipal inflows and compare them against the observed municipal receipts reported in Banco de México's series CE166.

Figure 4.2 plots the Pearson correlation between our estimates aggregated at the municipal level and the official inflows for each year. The results report correlation coefficients consistently between 0.75 and 0.85, suggesting that our estimation procedure successfully captures a significant portion of the variation in remittance inflows across Mexican municipalities. While this result validates our baseline approach, we acknowledge that it does not capture all the heterogeneity in remittance flows, which is expected given the simplifying assumptions we made.




Figure 4.2: Correlation of estimated municipal remittance inflows with official data. Each point represents the Pearson correlation coefficient between the estimated municipal inflows (aggregated from our bilateral estimates) and the actual municipal inflows recorded in Banxico's series CE166 for a given year.

Source: Authors' own elaboration.

5. Empirical Methodology and Results

5.1. Methodology

We use the following Poisson Psuedo Maximum Likelihood event-study specification to analyze the impact of the Hurricane Ian shock on remittance inflows to Mexican municipalities:

$$\mathbb{E}[R_{mt} \mid X_{mt}] = \exp(\alpha_t + \gamma_{st} + \sum_{k \neq -1} \beta_k \cdot w_{m,FL} \cdot \mathbb{1}\{t - T_0 = k\} + \epsilon_{st}) \quad (5.1)$$

This methodology is motivated by Silva and Tenreyro, 2006, who show that the Poisson Pseudo Maximum Likelihood (PPML) estimator is consistent in the presence of heteroskedasticity and excess zeros, which are common features of remittance data at the municipality level. The event-study specification allows us to estimate dynamic treatment effects of the shock over time, while controlling for time-varying unobservables at the state level through γ_{st} and common shocks through α_t . The interaction term $w_{m,FL} \cdot \mathbb{1}\{t - T_0 = k\}$ captures the differential effect of the shock on municipalities with varying levels of exposure to Florida, where $w_{m,FL}$ is a measure of pre-shock exposure to Florida and $\mathbb{1}\{t - T_0 = k\}$ is an indicator for each time period relative to the shock. The coefficients β_k thus represent the change in remittance inflows for a one-unit increase in Florida exposure at each time period k relative to the shock. We cluster standard errors at the municipality level to account for potential correlation within municipalities over time. Note that as a robustness check, we also consider clusters at the spatial level, using 10km, 25km, 50km, and 100km-radius clusters to account for potential spatial correlation in remittance responses among nearby municipalities. See the online appendix for results with spatial clustering.

5.2. Robustness checks

miguel TDL: insert charts (or just ramsay RESET test) to check for conditional mean form misspecification. can put additional stuff in online appendix.

5.3. Results

insert PPML event study results here.

insert interpretation and discussion of results here.

6. Investigating Spillover Effects

A fundamental requirement for standard causal inference models is the Stable Unit Treatment Value Assumption (SUTVA), formalized by Rubin (1980). SUTVA assumes that the potential outcomes of any unit depend exclusively on its own treatment assignment, excluding the possibility of any indirect spillover between units.

In our context, U.S. states not directly exposed to a local shock might still be affected if control units are exposed to the treatment through indirect channels, leading to biased estimates (Clarke, 2017). For instance, given that California and Texas all host a big proportion of Mexican migrants, we hypothesize that these two states might have played a role in offsetting the negative Florida shock. If a Mexican municipality highly dependent on Florida remittances experiences an income shock due to Hurricane Ian, connected migrants residing in Texas might respond by increasing their transfers to compensate for the drop. If this was the case, the estimated Florida-exposure effect on total remittances would capture the net effect of the shock after within-network adjustment, rather than the direct fall in Florida-origin remittances.

Therefore, we argue that our estimates are subject to general equilibrium effects. We base our analysis on the exposure-mapping logic of Aronow and Samii (2017) to define indirect exposure through alternative migrant networks. Rather than assuming that a municipality's remittances respond only to Florida exposure, we allow the response to vary with exposure to other major U.S. destinations, especially Texas and California.

6.1. Empirical Strategy

The biggest hosts of Mexican migrants are Texas and California (include data here); thus, it is likely that those municipalities exposed to the Florida shock might also be connected to a certain extent to California and Texas. Plot ?? helps us visualize that there exists indeed some migrant networks spanning across Florida and California and Florida and Texas.

Figure 6.1 suggests that there has been a dip in all 3 states in 2022Q4, after the Hurricane Ian hit Florida. However, we are interested in examining how the inflow of remittances from Texas and California evolved around the shock relative to the evolution of Florida inflows, specifically for the municipalities that are jointly highly exposed to Florida and Texas, as there exists parallel trends in the remittance patterns.




Figure 6.1: Evolution of Remittances Inflows from FL, TX and CA to all Mexican municipalities

Let's define a measure of joint exposure FL-TX of Mexican municipality m as the geometric mean of its exposure to Florida and its exposure to Texas, as follows:

$$JointExposure_m^{FL,TX} = \text{sqrt}(Exposure_m^{FL} \times Exposure_m^{TX})$$

Similarly, we can also define a measure of joint exposure FL-CA.

Note that in the TX-FL regression, the sample is restricted to the highly exposed municipalities (As defined above). Similarly, in the CA-FL regression, the sample is also restricted to the corresponding high jointly exposed municipalities.

We include α_t as the quarter FE, γ_{ms} as the corridor MX municipality - US State FE and $Post_t$ is defined such that $Post_t = 1$ for quarters from 2022Q4 onward, and 0 otherwise. Since Florida is the baseline category, the coefficient of interest β_{TX} (or β_{CA}) compares the post-2022Q4 change in Texas remittances to the post-2022Q4 change in Florida remittances (among our subsamples of interest). To ensure robustness of results, we also clustered at corridor and quarter level.

In other words, the coefficients β_{TX} and β_{CA} capture respectively:

$$\begin{aligned}\beta_{TX} &= [\Delta \log R_{m,Texas}] - [\Delta \log R_{m,Florida}] \\ \beta_{CA} &= [\Delta \log R_{m,California}] - [\Delta \log R_{m,Florida}]\end{aligned}$$

6.2. Results




Figure 6.2: Change in log remittances inflows from TX and CA rel. to FL to high joint exposure municipalities

7. Conclusion

7.1. Summary of Findings

7.2. Policy Implications

7.3. Limitations and Future Research

Bibliography

- Albert, C., & Monras, J. (2022). Immigration and spatial equilibrium: The role of expenditures in the country of origin. *American Economic Review*, 112(11), 3763–3802. <https://doi.org/10.1257/aer.20211241>
- Alcaraz, C., Chiquiar, D., & Salcedo, A. (2012). Remittances, schooling, and child labor in Mexico. *Journal of Development Economics*, 97(1), 156–165. <https://doi.org/10.1016/j.jdevec.2010.11.004>
- Amuedo-Dorantes, C., & Pozo, S. (2006). Remittances as insurance: Evidence from mexican immigrants. *Journal of Population Economics*, 19, 227–254. <https://doi.org/10.1007/s00148-006-0079-6>
- Aronow, P. M., & Samii, C. (2017). Estimating average causal effects under general interference, with application to a social network experiment. *The Annals of Applied Statistics*, 11(4), 1912–1947. <https://doi.org/10.1214/16-AOAS1005>
- Banco de México. (2026a). *Estado de origen de los ingresos por remesas provenientes de Estados Unidos (CE168)* [Sistema de Información Económica (SIE), Banco de México]. Retrieved May 26, 2026, from <https://www.banxico.org.mx/SieInternet/consultarDirectorioInternetAction.do?accion=consultarCuadro&idCuadro=CE168&locale=es>
- Banco de México. (2026b). *Ingresos por remesas familiares (CA11)* [Sistema de Información Económica (SIE), Banco de México]. Retrieved May 26, 2026, from <https://www.banxico.org.mx/SieInternet/consultarDirectorioInternetAction.do?accion=consultarCuadroAnalitico&idCuadro=CA11§or=1&locale=es>
- Banco de México. (2026c). *Ingresos por remesas, distribución por municipio (CE166)* [Sistema de Información Económica (SIE), Banco de México]. Retrieved May 26, 2026, from <https://www.banxico.org.mx/SieInternet/consultarDirectorioInternetAction.do?accion=consultarCuadro&idCuadro=CE166&locale=es>
- Bettin, G., Jallow, A., & Zazzaro, A. (2025). Responding to natural disasters: What do monthly remittance data tell us? *Journal of Development Economics*, 174, 103413. <https://doi.org/10.1016/j.jdevec.2024.103413>
- Boustan, L. P., Kahn, M. E., Rhode, P. W., & Yanguas, M. L. (2020). The effect of natural disasters on economic activity in US counties: A century of data. *Journal of Urban Economics*, 118, 103257. <https://doi.org/10.1016/j.jue.2020.103257>

- Bucci, L., Alaka, L., Hagen, A., Delgado, S., & Beven, J. (2026, March 12). *Tropical cyclone report: Hurricane Ian (AL092022), 23–30 September 2022* (tech. rep.). National Hurricane Center, National Oceanic and Atmospheric Administration. Retrieved May 26, 2026, from http://www.nhc.noaa.gov/data/tcr/AL092022_Ian.pdf
- Butts, K. (2023). Difference-in-differences estimation with spatial spillovers. <https://arxiv.org/abs/2105.03737>
- Caballero, M. E., Cadena, B. C., & Kovak, B. K. (2018). Measuring geographic migration patterns using Matrículas Consulares. *Demography*, 55(3), 1119–1145. <https://doi.org/10.1007/s13524-018-0675-6>
- Caballero, M. E., Cadena, B. C., & Kovak, B. K. (2023). The international transmission of local economic shocks through migrant networks. *Journal of International Economics*, 145, 103832. <https://doi.org/10.1016/j.jinteco.2023.103832>
- Card, D. (2001). Immigrant inflows, native outflows, and the local labor market impacts of higher immigration. *Journal of Labor Economics*, 19(1), 22–64. <https://doi.org/10.1086/209979>
- Clarke, D. (2017). *Estimating difference-in-differences in the presence of spillovers* (MPRA Paper No. 81604). University Library of Munich. <https://mpra.ub.uni-muenchen.de/81604/>
- Hudgens, M. G., & Halloran, M. E. (2008). Toward causal inference with interference. *Journal of the American Statistical Association*, 103(482), 832–842. <https://doi.org/10.1198/016214508000000292>
- Instituto de los Mexicanos en el Exterior. (2026). *Estadísticas: Información estadística de la población mexicana en el exterior*. Secretaría de Relaciones Exteriores. Retrieved April 1, 2026, from <https://ime.gob.mx/estadisticas>
- Melitz, M. J. (2003). The impact of trade on intra-industry reallocations and aggregate industry productivity. *Econometrica*, 71(6), 1695–1725. <https://doi.org/10.1111/1468-0262.00467>
- Mora-Rivera, J., & van Gameren, E. (2021). The impact of remittances on food insecurity: Evidence from Mexico. *World Development*, 140, 105349. <https://doi.org/10.1016/j.worlddev.2020.105349>
- Munshi, K. (2003). Networks in the modern economy: Mexican migrants in the U.S. labor market. *The Quarterly Journal of Economics*, 118(2), 549–599. <https://doi.org/10.1162/003355303321675455>

- Rubin, D. B. (1980). Randomization analysis of experimental data: The Fisher randomization test comment. *Journal of the American Statistical Association*, 75(371), 591–593. <https://doi.org/10.1080/01621459.1980.10477517>
- Silva, J. M. C. S., & Tenreiro, S. (2006). The log of gravity. *The Review of Economics and Statistics*, 88(4), 641–658. <https://doi.org/10.1162/rest.88.4.641>
- Taylor, J. E., Mora, J., Adams, R. H., & Lopez-Feldman, A. (2005). *Remittances, inequality and poverty: Evidence from rural Mexico* (Working Paper No. 05-003). University of California, Davis, Department of Agricultural and Resource Economics. <https://doi.org/10.22004/ag.econ.60287>
- World Bank. (2024). *Personal remittances, received (current US\$)* [World Development Indicators, series BX.TRF.PWKR.CD.DT]. Retrieved May 26, 2026, from <https://data.worldbank.org/indicator/BX.TRF.PWKR.CD.DT?locations=MX>
- Yang, D., & Choi, H. (2007). Are remittances insurance? evidence from rainfall shocks in the Philippines. *The World Bank Economic Review*, 21(2), 219–248. <https://doi.org/10.1093/wber/lhm003>